

Yandex

Yandex

Yet another password hashing talk

Evgeny Sidorov

Make out my accent

- › Yescript
- › Scrypt
- › Bcrypt

Agenda

1

INTRO

2

Algorithms

3

Protocol designs

4

Conclusions

Introduction



Intro

- › Surprisingly passwords are still alive
- › Passwords will probably always exist
- › “Something that you know” in 2FA

Password Hashing

- › Non-interactive operations (password based encryption etc.)
- › Interactive / online operations (password based authentication in a web service etc.)
- › This talk is about interactive / online operations

Attacker and Defender

- › Defender uses traditional multicore server CPUs (bigger cache per core, bigger RAM available)
- › Attacker can use GPUs, FPGAs, ASICs
- › GPU is the most important case as they're ready to use

Algorithms



The role of PHA (interactive mode)

- › Itself not the perfect security measure
- › Must be considered as a “last hope”
- › Must give users time to change their passwords in case of DB compromise (total compromise)

Strategies behind PHAs

- › Optimise as much as possible for CPUs, their capabilities and RAM access patterns
- › Make GPUs, ASICs, FPGAs ineffective/expensive by using more RAM (and other of their bottlenecks)
- › Memory hard - trading memory for computations leads to much bigger time and computational costs

Yescrypt



Yescrypt

- › Designed by Solar Designer (aka Alexander Peslyak)
- › Got special recognition from PHC
- › Optimised for “password hashing at scale”

Yescrypt

- › Fixes scrypt's TMTO
- › Uses not only random reads but also random writes (increases the amount of recomputation)
- › Replaces Salsa20/8 with custom "pwxform" to achieve bcrypt-like resistance to GPUs
- › Some tweaks in parallelism (moved to deeper level)

Script

PBKDF2-HMAC-SHA-256

SMix

BlockMix

Salsa20/8

PBKDF2-HMAC-SHA-256

pxform

- › An optimisation for low memory case
- › Uses 12kb s-boxes (8kb active lookups, 16 byte individual lookup) and it's friendly for CPU L1 cache
- › Uses integer multiplications (higher latency on FPGAs/ASICs)
- › Friendly to SIMD instructions

Pros&Cons

- › Optimised for low memory case (bcrypt-like resistance)
- › Fixes Script drawbacks
- › Pretty complex
- › There is no resistance to side channel attacks


Argon2

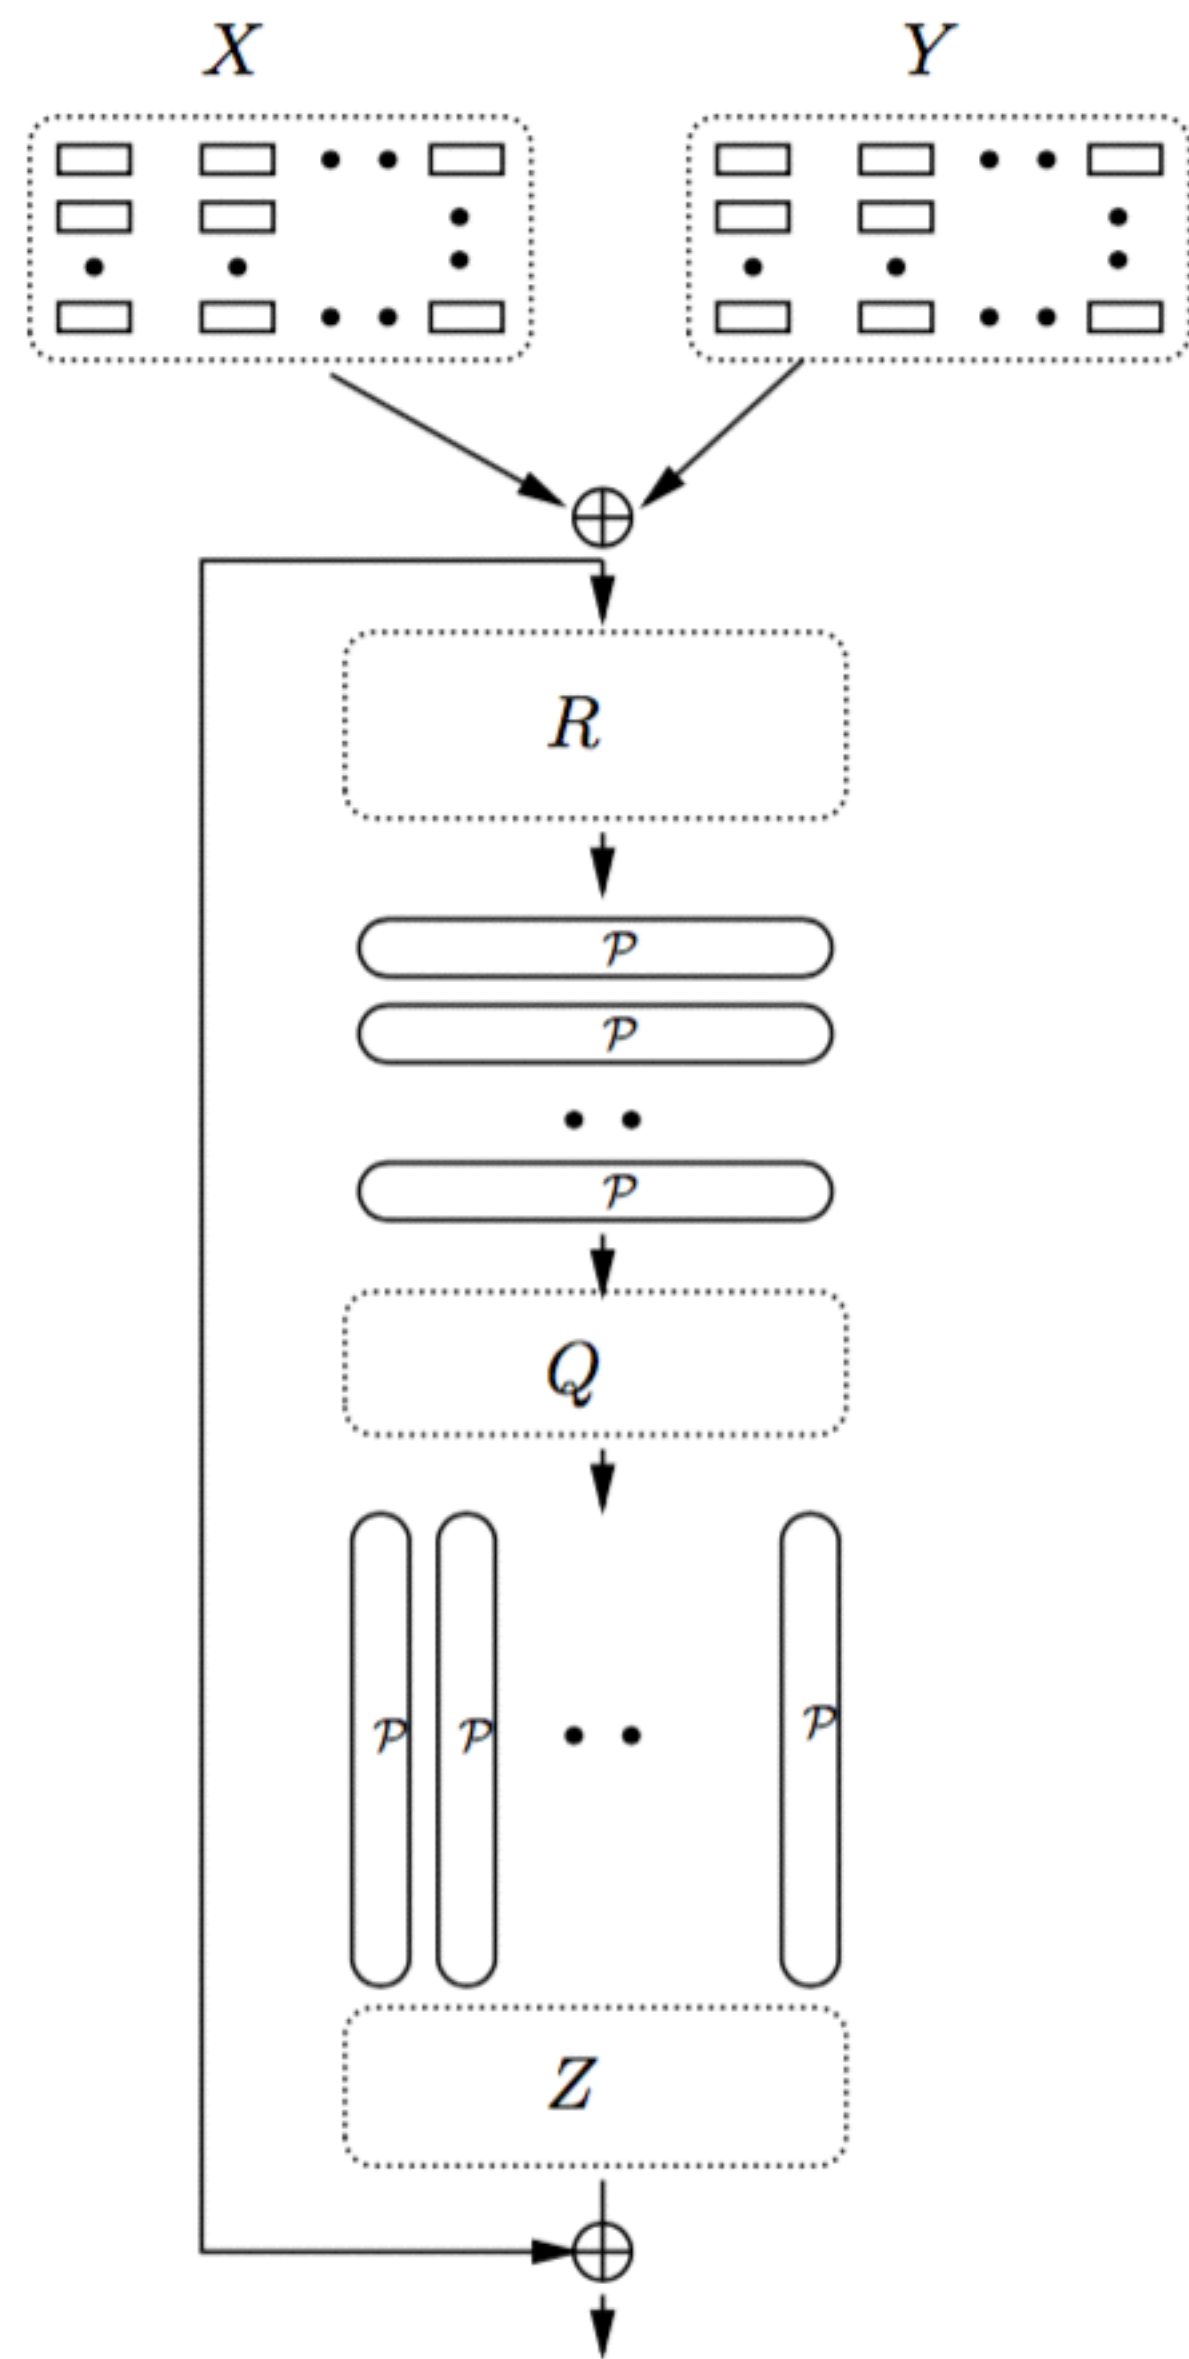


Argon2

- › Simple, tuneable & secure
- › Three versions: i, d, id

Argon2





Compression function

$a \leftarrow a + b$
 $d \leftarrow (d \oplus a) \ggg 32$
 $c \leftarrow c + d$
 $b \leftarrow (b \oplus c) \ggg 24$
 $a \leftarrow a + b$
 $d \leftarrow (d \oplus a) \ggg 16$
 $c \leftarrow c + d$
 $b \leftarrow (b \oplus c) \ggg 63$

(a) Blake2 G function.

$a \leftarrow a + b + 2 \cdot \text{lsb}(a) \cdot \text{lsb}(b)$
 $d \leftarrow (d \oplus a) \ggg 32$
 $c \leftarrow c + d + 2 \cdot \text{lsb}(c) \cdot \text{lsb}(d)$
 $b \leftarrow (b \oplus c) \ggg 24$
 $a \leftarrow a + b + 2 \cdot \text{lsb}(a) \cdot \text{lsb}(b)$
 $d \leftarrow (d \oplus a) \ggg 16$
 $c \leftarrow c + d + 2 \cdot \text{lsb}(c) \cdot \text{lsb}(d)$
 $b \leftarrow (b \oplus c) \ggg 63$

(b) BlaMka G function.

BlaMka function

<https://eprint.iacr.org/2015/136.pdf>

Argon2 design principles

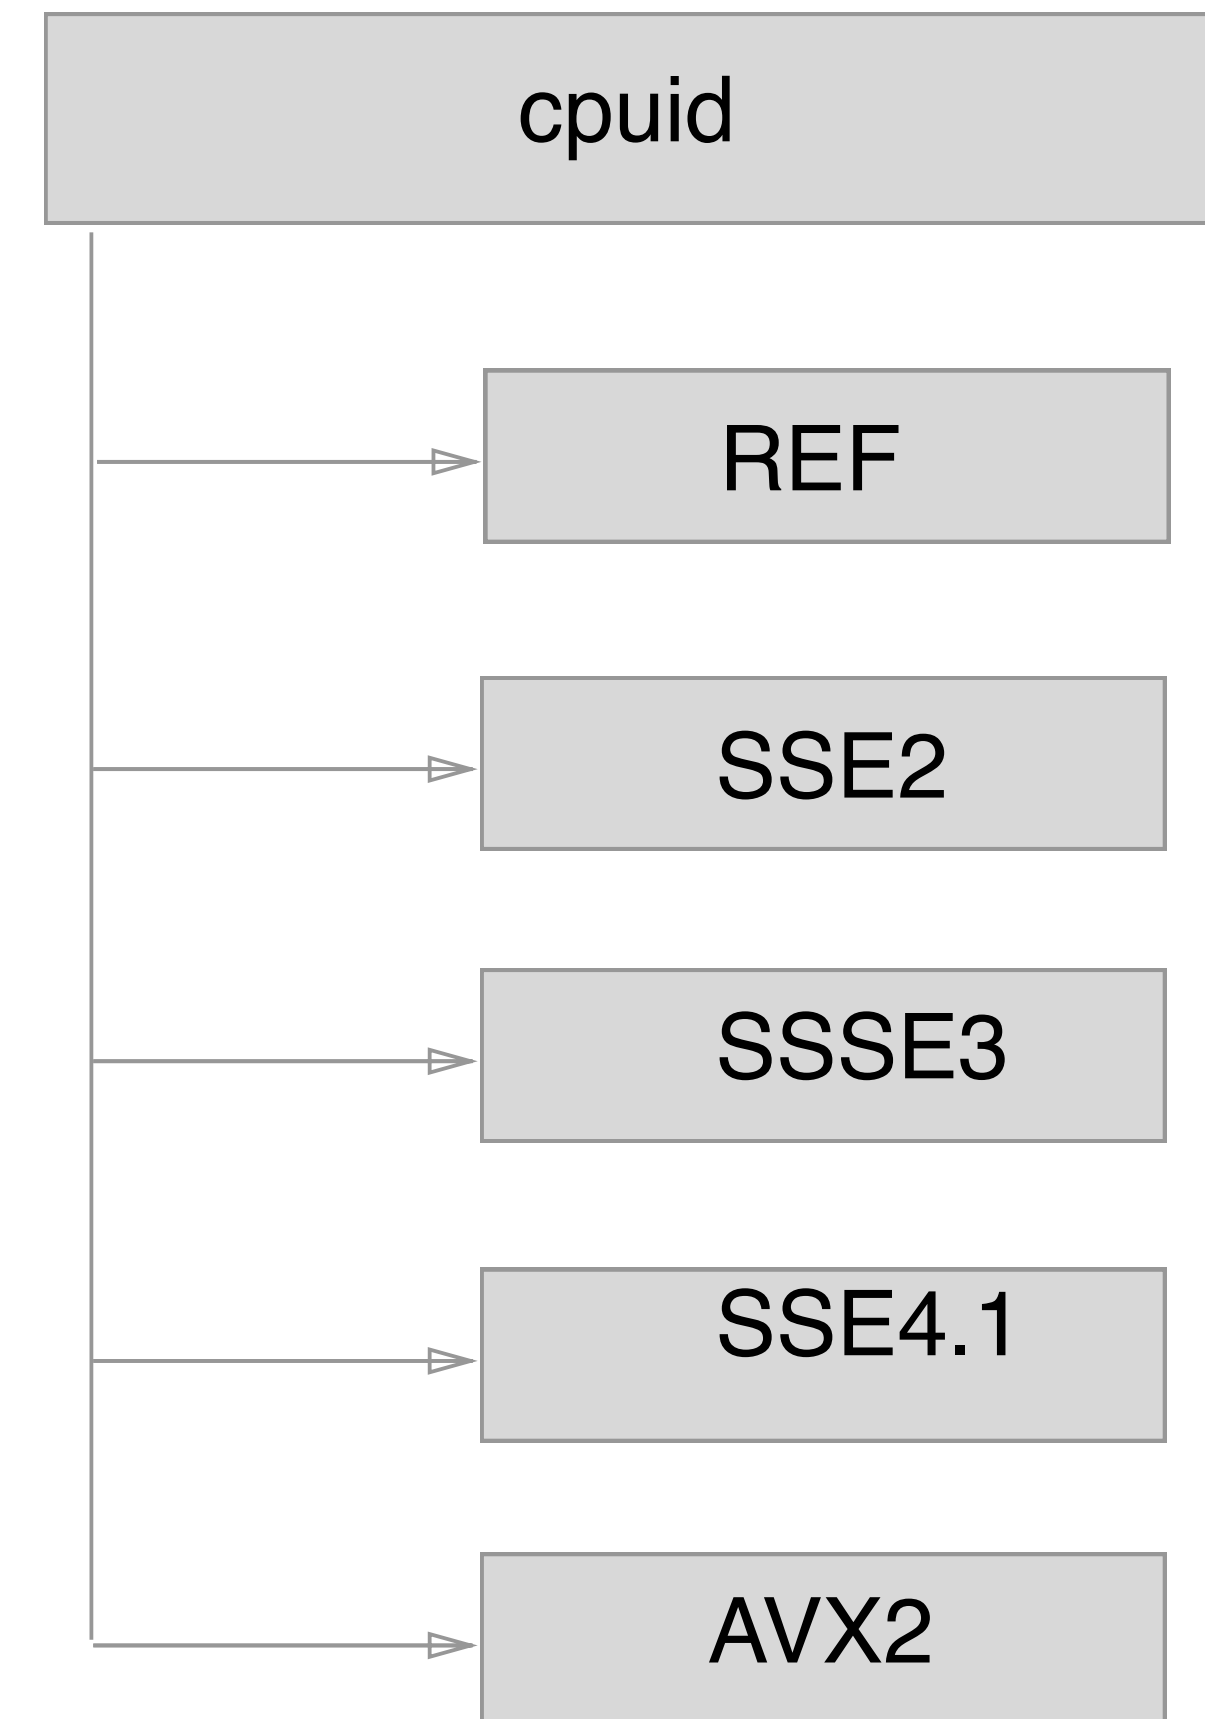
- › Operates on 1kb memory blocks
- › Optimised for x64, SIMD friendly
- › Uses Blake2B and BlaMka as underlying primitives
- › BlaMka increases the circuit depth and makes ASIC implementation more complex (uses integer multiplication)

PHC implementation

- › PHC implementation of Argon2 uses “-march=native”
- › Can easily run up against “Illegal instructions exception”
- › Compatibility costs 15-20% of performance

Our implementation

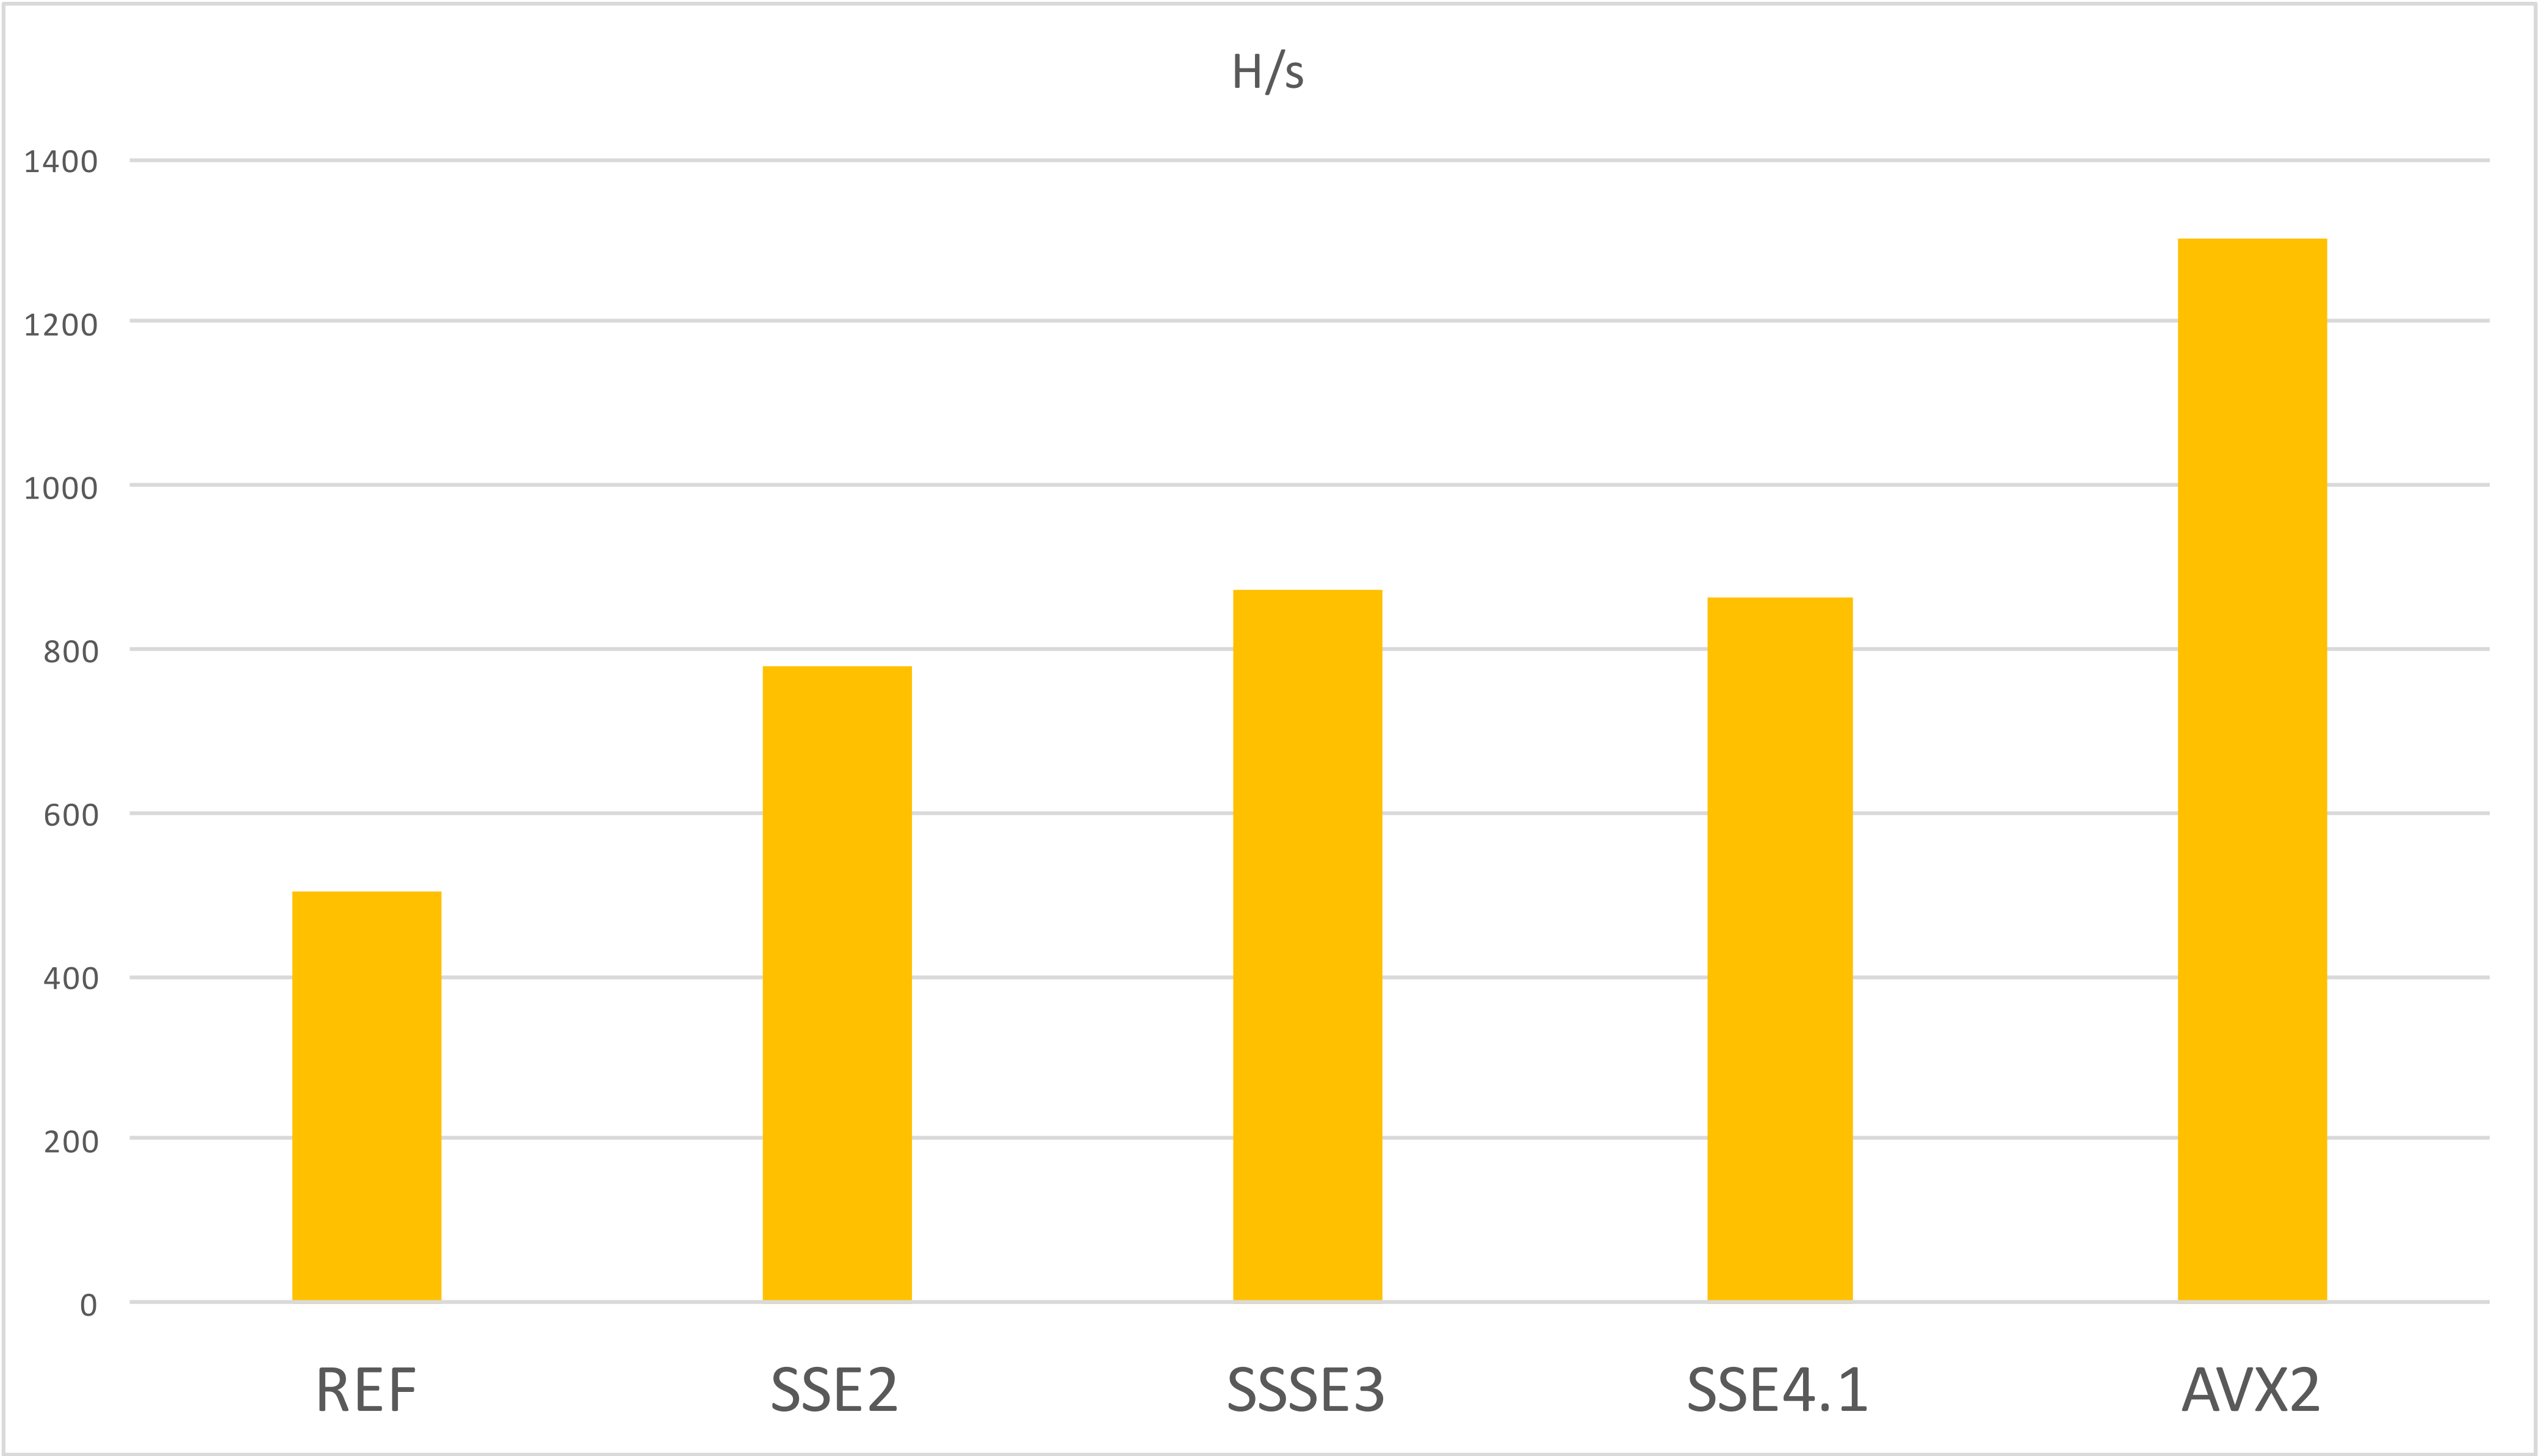
- › Employs SIMD as much as possible
- › Runtime CPU dispatching to run the version optimised for current CPU
- › Makes deployment easier (especially for clouds)
- › SIMD-optimized Blake2B



Naming

Argonishche = Argon + itch (-ищ)

Benchmark Argon2d (1, 2048, 1) c4.8xlarge



Argonishche (Argon2 implementation)

- › <https://github.com/yandex/argon2>
- › C++14 (constexpr, partial templates etc.)
- › BSD 3-clause license
- › AVX512 (Intel Skylake) is in plans

Protocol Designs



Motivation

- › Algorithms are not enough
- › Large parameters are hardly affordable for high-loaded services
- › Offline attacks are possible
- › Need to apply protocol level mitigations

Protocol designs

- › Local parameters
- › Big ROMs (“security through obesity”)
- › Crypto Anchors
- › Partially Obvious PRFs

Local parameters

- › Add secret key to password hashing algorithm
- › Secret key stored in the application conf, not in the DB
- › SQL injection or just access to production DB become insufficient
- › RCE in the application is needed

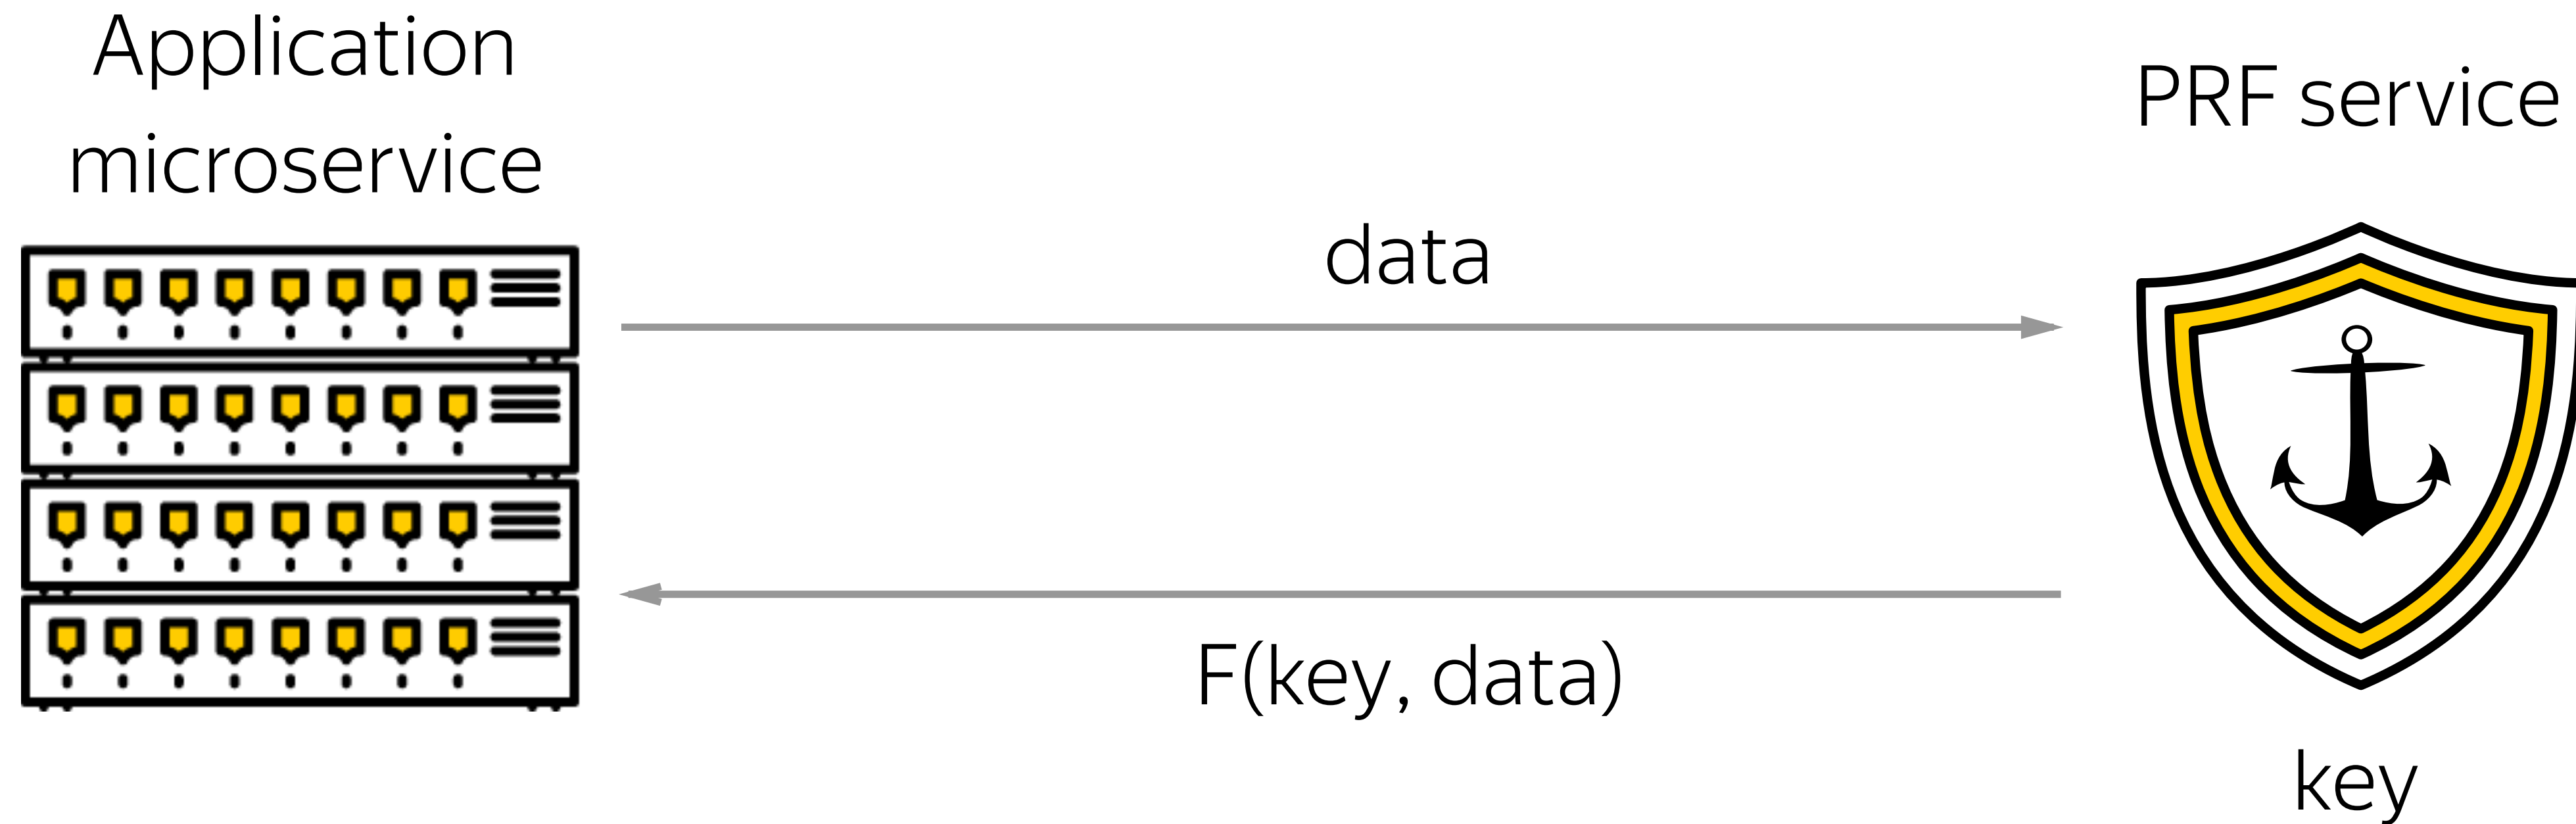
Big ROM pros

- › A relatively big pre-filled shared memory area
- › using many cores more expensive for the shared ROM, less cores are good (like in a CPU)
- › For example, 112Gb ROM is harder to steal

Big ROM cons

- › Make it harder to switch to a new PHA (you'll have to live with the ROM)
- › ROM-as-a-service amortisation approach (pieces of ROM can be scattered around less powerful servers, network is used to query values)
- › Leads to additional steps in case of server reboot etc.

Crypto Anchors



Crypto Anchors

- › Tiny PRF services (can apply HMAC for example)
- › Turn offline attacks into online attacks (attacker just can't go home and have fun with hashes)
- › In case of HMAC key updates are impossible (need to add one more layer)

PASS: PRF service

PASS: Strengthening and Democratizing Enterprise Password Hardening

Ari Juels

Jacobs Technion-Cornell Institute

Cornell Tech

with D. Akhawe (Dropbox), A. Athalye (MIT), R. Chatterjee (Cornell), A. Everspaugh (UWisc), T. Ristenpart (Cornell Tech), S. Scott (Royal Holloway)

$$e: G_1 \times G_2 \rightarrow G_T$$


$$e(a^u, b^v) = e(a, b)^{uv}$$

Bilinear pairing


Blinding $x := H(P)^r$

t, x

$y := F_k(t, x) = e(H(t), x)^k$



y



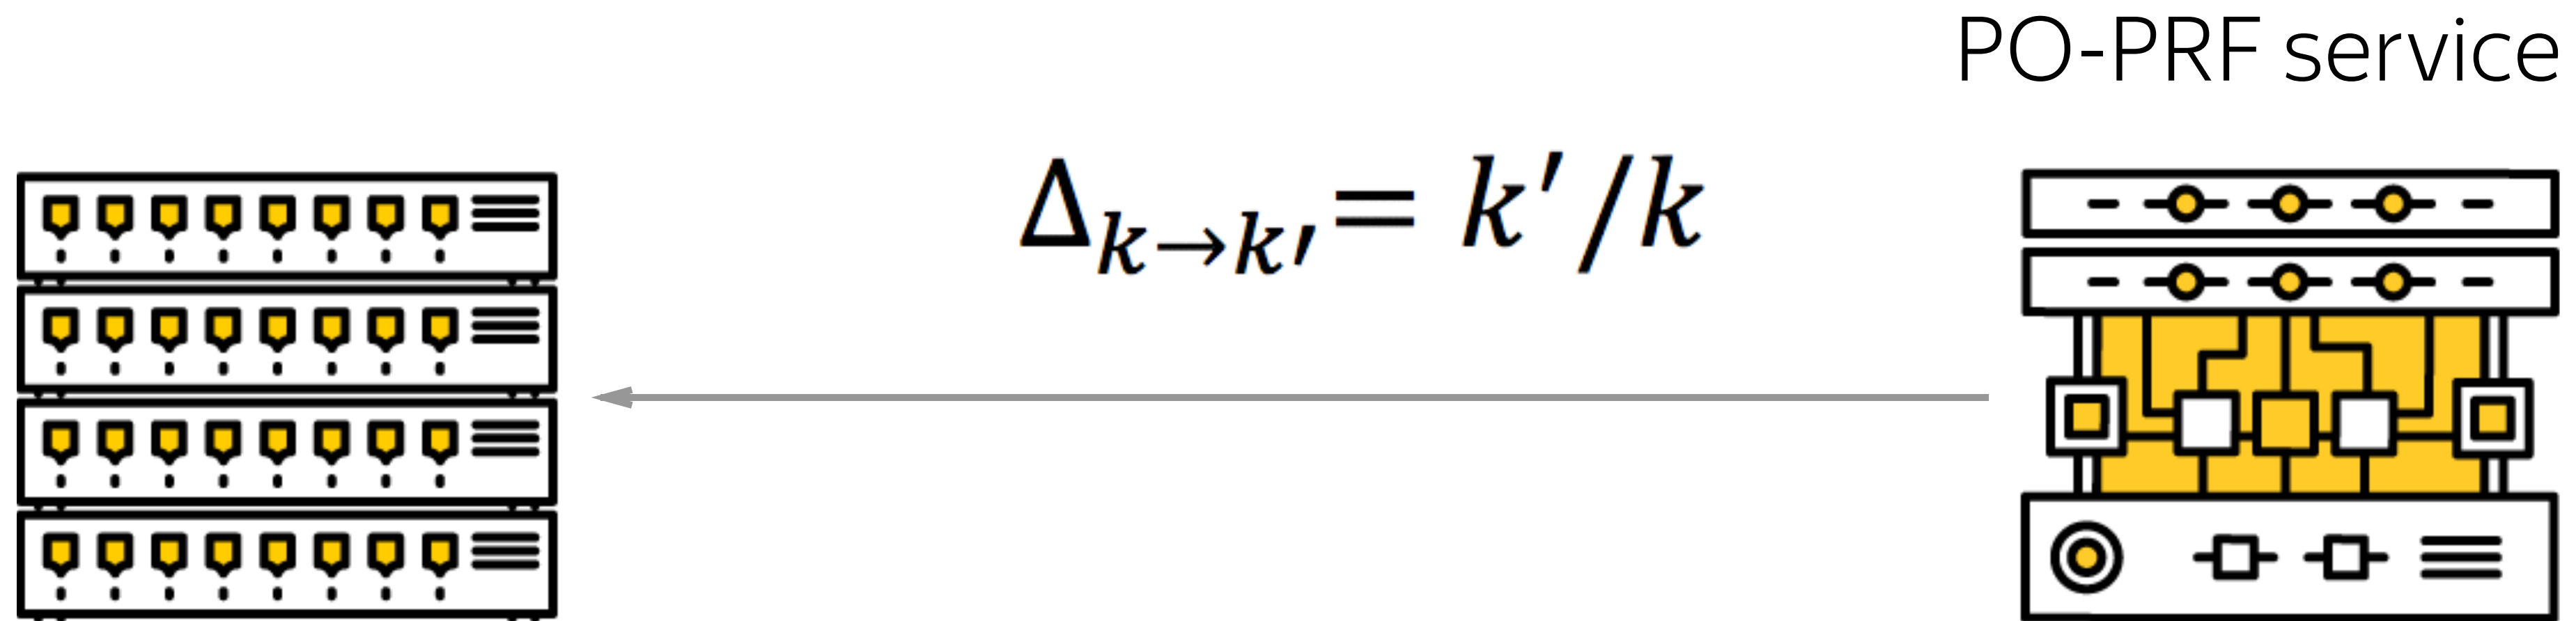
PO-PRF service

$$z := y^{\frac{1}{r}} = e(H(t), H(p))^{k \cdot r \cdot \frac{1}{r}}$$

$$= e(H(t), H(p))^k$$

Unblinding

Key rotation



$$z' := z^{\frac{k'}{k}} = e(H(t), H(p))^{k * \frac{k'}{k}} = e(H(t), H(p))^{k'}$$

Pythia PRF

- › The idea is like those in PASS
- › Uses Zero Knowledge Proof scheme to proof correctness of calculations
- › Multitenancy

Conclusions



Conclusions

- › Good algorithms are not enough for online password hashing
- › There are some good approaches to make up for it
- › Offline attacks can be turned into online attacks

Acknowledgements

- › Dmitry Khovratovich for useful discussions about Yescrypt and Argon2
- › Alexander Peslyak for discussions about Yescrypt
- › Igor Klevanets for reviewing the Argon2 library (Argonische)

Useful links

- › <https://eprint.iacr.org/2015/644> Pythia PRF Service
- › <https://github.com/yandex/argon2> Yandex's Argon2 impl
- › <http://bit.ly/2vS0I7B> Pythia ref. implementation

Questions?

Evgeny Sidorov

Information Security Engineer



e-sidorov@yandex-team.ru

<https://github.com/yandex/argon2>